



SINGLE PHOTON DOUBLE SLIT SIMULATION

Seminar Talk:


Parallelisation of Physics Calculations on GPUs with Cuda

by **Xianrui Yin**



CONTENTS



- Double slit experiment (*wave v.s. particle*)
- Simulation based on physics argument
- Parallelisation on GPU
- Why it works?
- What else could be done?




DOUBLE SLIT EXPERIMENT

[1]


SINGLE PHOTON



SET UP THE LAB



The photon follows a straight line path
and hit one of the detectors on the screen



[3]

2M
..
..
..
..
..
..
..
2
1
0

- 
- Photons are oscillating

$$e = [\cos\varphi, \sin\varphi]$$

- Each detector has a polarization vector

$$p = [p_x, p_y]$$

- When photon lands on a detector, it changes its polarization

$$p_{i+1} = \gamma p_i + (1 - \gamma)e_i$$

- A photon is detected if the threshold is met

$$\Theta(p^2 - r) > 0$$

E-S Correspondence

Experiment	Simulation
line source	choose y and β randomly
photon phase change φ	$e = [\cos\varphi, \sin\varphi]$
polarization of the detector material	$p = [p_x, p_y]$
Interaction between photon & detector	$p_{i+1} = \gamma p_i + (1 - \gamma)e_i$
sensitivity	$p^2 > r, r \text{ is random}$



CPU CODE REVIEW



LET'S MOVE TO GPU!



HOW TO PARALLELIZE?

Can we do everything all together?

Sadly not...

The polarization of the detectors must be updated sequentially.

SPLIT THE PROCEDURE INTO 2+1 PARTS!

- `__global__ void sample(...)`
 - Sampling can be done concurrently
 - Let each thread generate some random numbers
 - Record the target detector index and the phase of the photon
-
- `__global__ void detect(...)`
 - All the detectors are independent!
 - Send the pre-generated photons one-by-one

SPLIT THE PROCEDURE INTO 2+1 PARTS!

- `__global__ void sample(...)`
 - Sampling can be done concurrently
 - Let each thread generate some random numbers
 - Record the target detector index and the phase of the photon
-how to communicate?
- `__global__ void detect(...)`
 - All the detectors are independent!
 - Send the pre-generated photons one-by-one

Event array

thread

0

$\phi_0^{(0)}$ $idx_0^{(0)}$ $\phi_1^{(0)}$ $idx_1^{(0)}$

1

$\phi_0^{(1)}$ $idx_0^{(1)}$ $\phi_1^{(1)}$ $idx_1^{(1)}$

2

⋮

⋮

⋮

⋮

⋮

⋮

⋮

← sample per thread →

detector

detector array

we don't know the
length of each row

0

1

2


R

(

(

1

(



- Sort<<<1, 1>>>(event, det, len, ...)

Set a large enough N

Use another 1d array to record
the number of photons in each
row


detector

0

1


2

⋮
⋮
⋮
⋮
⋮




SIMULATION RESULTS

- CPU



Photon aggregation 2^{12}




Photon aggregation: 2^{16}

Photon aggregation 2^{21}


CPU run time: **56.625009059906006**

- As we keep shooting photons




Photon aggregation 2^{25}

CPU run time: **793.3088808059692**



Let's do 2^{25} photons on GPU!



THANKS GOOGLE..

+-----+									
NVIDIA-SMI 460.32.03 Driver Version: 460.32.03 CUDA Version: 11.2									
+-----+									
GPU		Name		Persistence-M		Bus-Id		Disp.A	
Fan		Temp		Perf		Pwr:Usage/Cap		Memory-Usage	
								Volatile	
								Uncorr. ECC	
								GPU-Util	
								Compute M.	
								MIG M.	
+-----+									
0		Tesla T4		Off		00000000:00:04.0		Off	
N/A		44C		P8		9W / 70W		0MiB / 15109MiB	
								0%	
								Default	
								N/A	
+-----+									
+-----+									
Processes:									
GPU		GI		CI		PID		Type	
		ID		ID				Process name	
								GPU Memory	
								Usage	
+-----+									
No running processes found									
+-----+									

```
==170== NVPROF is profiling process 170, command: ./test60
==170== Profiling application: ./test60
==170== Profiling result:
Type      Time(%)      Time      Calls      Avg      Min      Max      Name
GPU activities:  91.78%  2.69613s      1  2.69613s  2.69613s  2.69613s  sort(float*, float*, int*, int, int, int)
                6.10%  179.17ms      1  179.17ms  179.17ms  179.17ms  sample(float*, int, float, float, int)
                2.12%  62.273ms      1  62.273ms  62.273ms  62.273ms  detect(float*, float*, int*, int*, int, int)
                0.00%  2.6560us      1  2.6560us  2.6560us  2.6560us  [CUDA memcpy DtoH]
API calls:  83.62%  2.93757s      1  2.93757s  2.93757s  2.93757s  cudaDeviceSynchronize
            16.26%  571.13ms      5  114.23ms  2.0700us  570.81ms  cudaMalloc
            0.08%  2.6445ms      4  661.13us  3.1760us  2.2373ms  cudaFree
            0.03%  1.0065ms      1  1.0065ms  1.0065ms  1.0065ms  cuDeviceGetPCIBusId
            0.01%  348.49us      1  348.49us  348.49us  348.49us  cuDeviceTotalMem
            0.00%  157.86us     101  1.5620us   142ns   67.217us  cuDeviceGetAttribute
            0.00%  51.012us      1  51.012us  51.012us  51.012us  cudaMemcpy
            0.00%  47.733us      3  15.911us  5.4950us  34.765us  cudaLaunchKernel
            0.00%  29.370us      1  29.370us  29.370us  29.370us  cuDeviceGetName
            0.00%  1.8070us      3    602ns   209ns   1.1570us  cuDeviceGetCount
            0.00%  1.3050us      2    652ns   185ns   1.1200us  cuDeviceGet
            0.00%    280ns      1    280ns   280ns    280ns  cuDeviceGetUuid
```



Let's look at some code!

WHY IT WORKS?

$$p_{i+1} = \gamma p_i + (1 - \gamma)e_i$$

- The sequence will converge if $\gamma \rightarrow 1^-$ to the mean of e_i if $\gamma \rightarrow 1^-$
- constructive interference: $\langle e_i \rangle$ is large
- destructive interference: $\langle e_i \rangle$ is small

$$\Theta(p^2 - r) > 0$$




MORE TO EXPLORE

- Transient behavior
-



Thank you for your attention!

- 
- [1] https://commons.wikimedia.org/wiki/File:Two-Slit_Experiment_Light.svg
 - [2] Dimitrova, Todorka & Weis, Antoine. (2008). The wave-particle duality of light: A demonstration experiment. American Journal of Physics. 76. 10.1119/1.2815364.
 - [3] Jin F. Towards a corpuscular model of optical phenomena. Groningen: s.n., 2011. 114 p.